

# NoSQL and MongoDB

NoSQL vs SQL, MongoDB



**SoftUni Team**  
**Technical Trainers**



**SoftUni**

**Software University**

<https://about.softuni.bg/>


# Table of Contents

- Relational and Non-Relational Database
- Database Scalability
- CAP Theorem
- Distributed Systems
- NoSQL database types
- MongoDB Overview
- CRUD Operations



[sli.do](https://sli.do)

**#csharp-db**



# Relational and Non-Relational Databases

Differences and Examples

# Relational Database

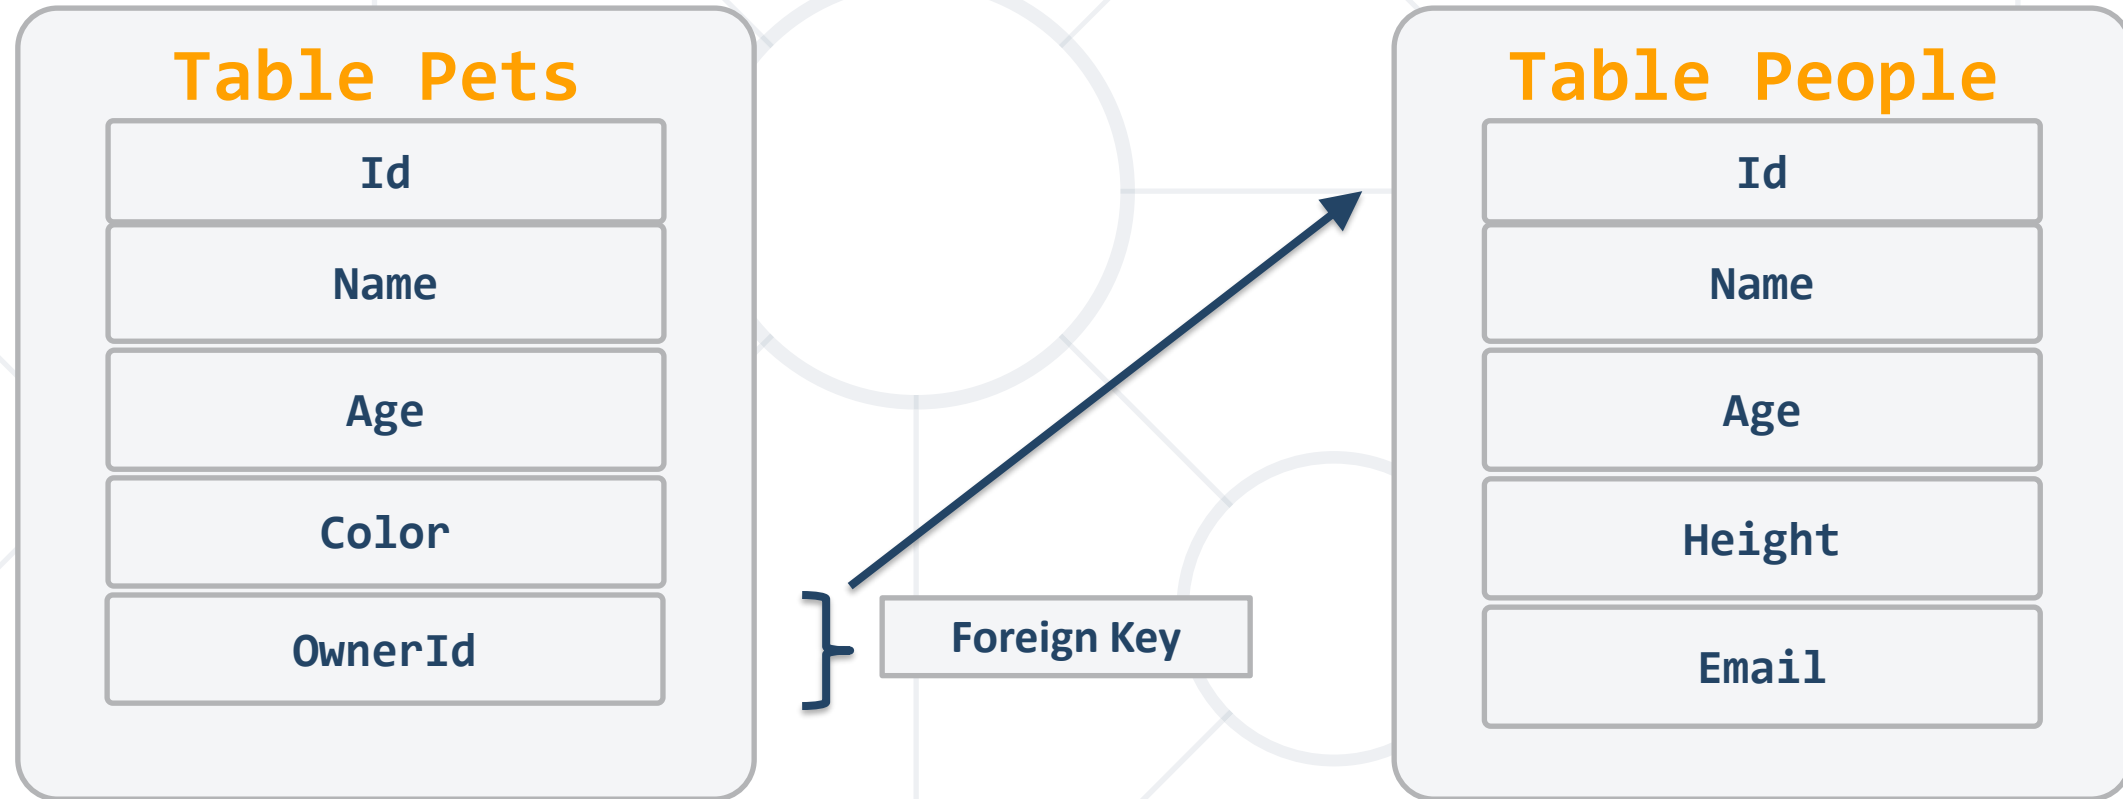
- Organizes data into one or more **tables** of **columns** and **rows**
- Unique **key** identifying each **row** of data
- Almost all relational databases use **SQL** to **extract** data



```
SELECT * FROM Students
```

- **Relations** between tables are done using **Foreign Keys (FK)**
- Such databases are **Oracle, MySQL, SQL Server**, etc..

# Relational Database - Example



# Non-relational Database (NoSQL)

- NoSQL Databases are non tabular, and store data differently than relational tables
- Key-value **stores**

```
{  
  "_id": ObjectId("59d3fe7ed81452db0933a871"),  
  "email": "peter@gmail.com",  
  "age": 22  
}
```

- **SQL** query is **not** used in NoSQL systems
- More **scalable** and **provide** superior **performance**



# Database Scalability



# Database Scalability (1)


## Vertical Scaling

( Increase size of instance (RAM , CPU etc.) )



## Horizontal Scaling

( Add more instances )



- The ability of a system's database to **scale up or down**, depending on the requirements
  - Enables the database to grow to a larger size to support more transactions
  - There are two types of database scalability
    - Vertical Scaling or Scale-up
    - Horizontal Scaling or Scale-out

- Refers to the process of adding more physical resources to the existing database server for improving the performance such as
  - Storage
  - Memory
  - CPU
- Helps in upgrading the capacity of the existing database server

# Vertical Scaling Pros and Cons

## ■ Pros

- It consumes less power
- You need to handle and manage just one system
- Cooling costs are less than horizontal scaling
- Implementation isn't difficult

## ■ Cons

- Risk of hardware failure which can cause bigger outages
- Limited scope of upgradeability in the future



- Adds more servers with less RAM and processors
  - The ability to **increase the capacity** by connecting multiple software or hardware entities in a such manner that they function as a single logical unit
  - If a cluster requires more resources to improve its performance and provide high availability, then **the administrator can scale-out by adding more servers** to the cluster

# Horizontal Scaling Pros and Cons


## ■ Pros

- Easy to upgrade
- Resilience is improved due to the presence of discrete, multiple systems
- Supports linear increases in capacity

## ■ Cons

- It has a bigger footprint in the Data Center
- Adds complexity to the system
- Introduces data syncing problems
- Dependent on the CAP theorem

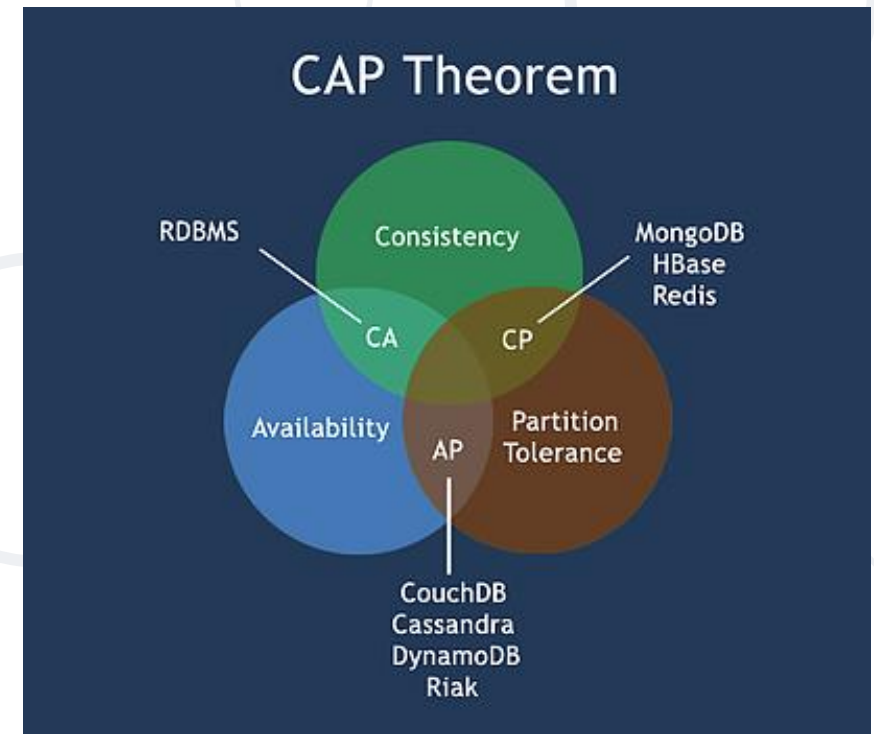




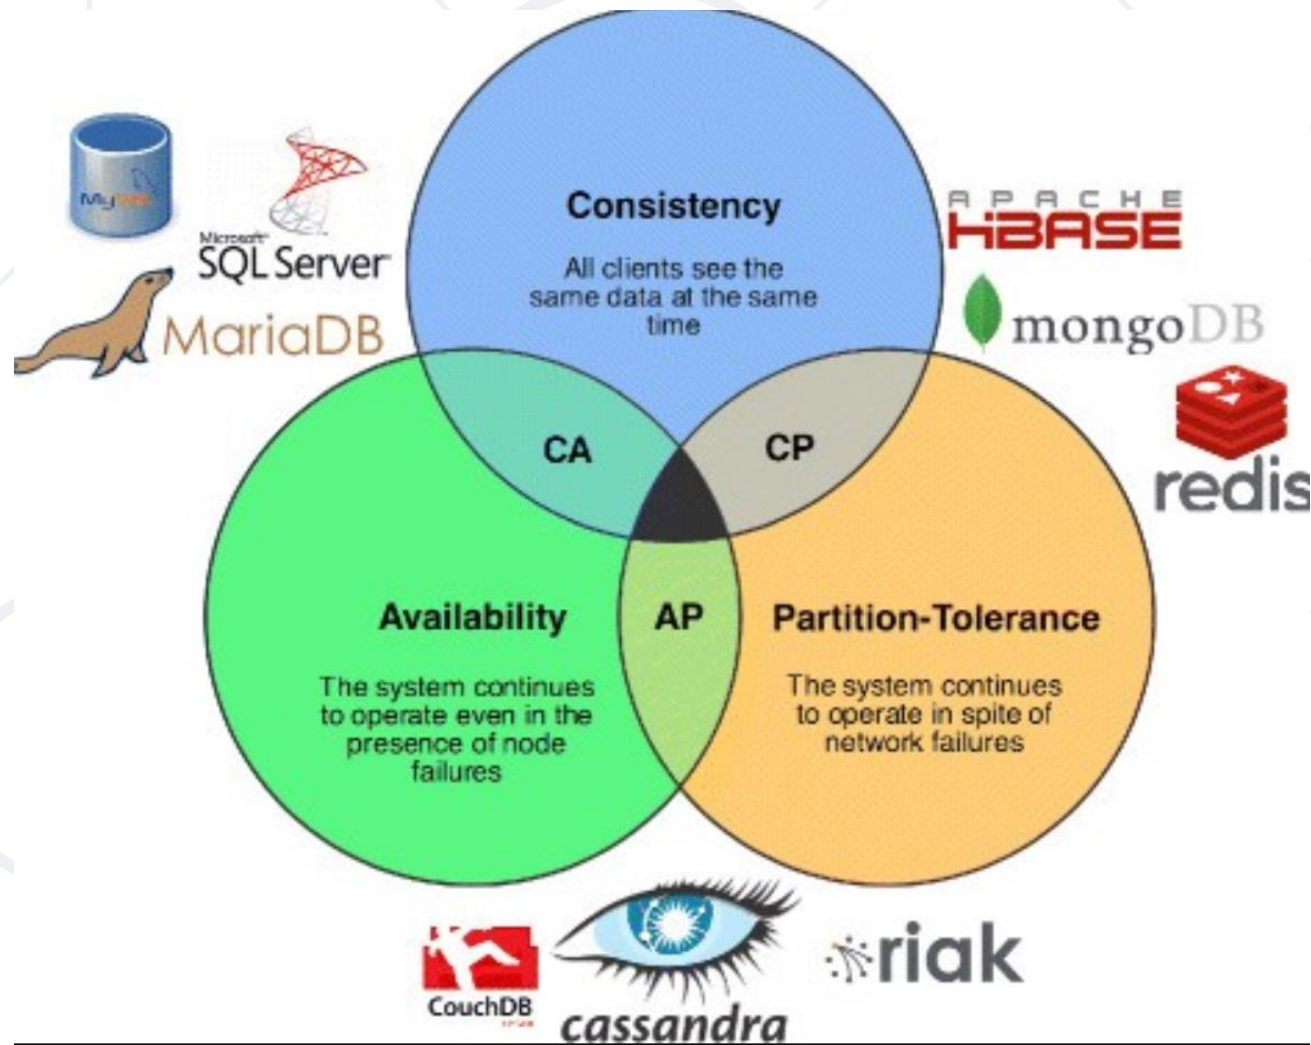
# CAP Theorem

# What is the CAP Theorem?

- The CAP theorem states that a distributed system can deliver **only two of three** desired characteristics
  - Consistency
  - Availability
  - Partition tolerance







- Consistency
  - All clients see the same data at the same time, no matter which node they connect to
    - Whenever data is written to one node, it must be instantly forwarded or replicated to all the other nodes in the system before the write is deemed 'successful'

# The 'CAP' in the CAP Theorem, Explained (2)


- Availability
  - Any client making a request for data gets a response, even if one or more nodes are down
- Partition tolerance
  - The cluster must continue to work despite any number of communication breakdowns between nodes in the system



# Distributed Systems

# Distributed Systems

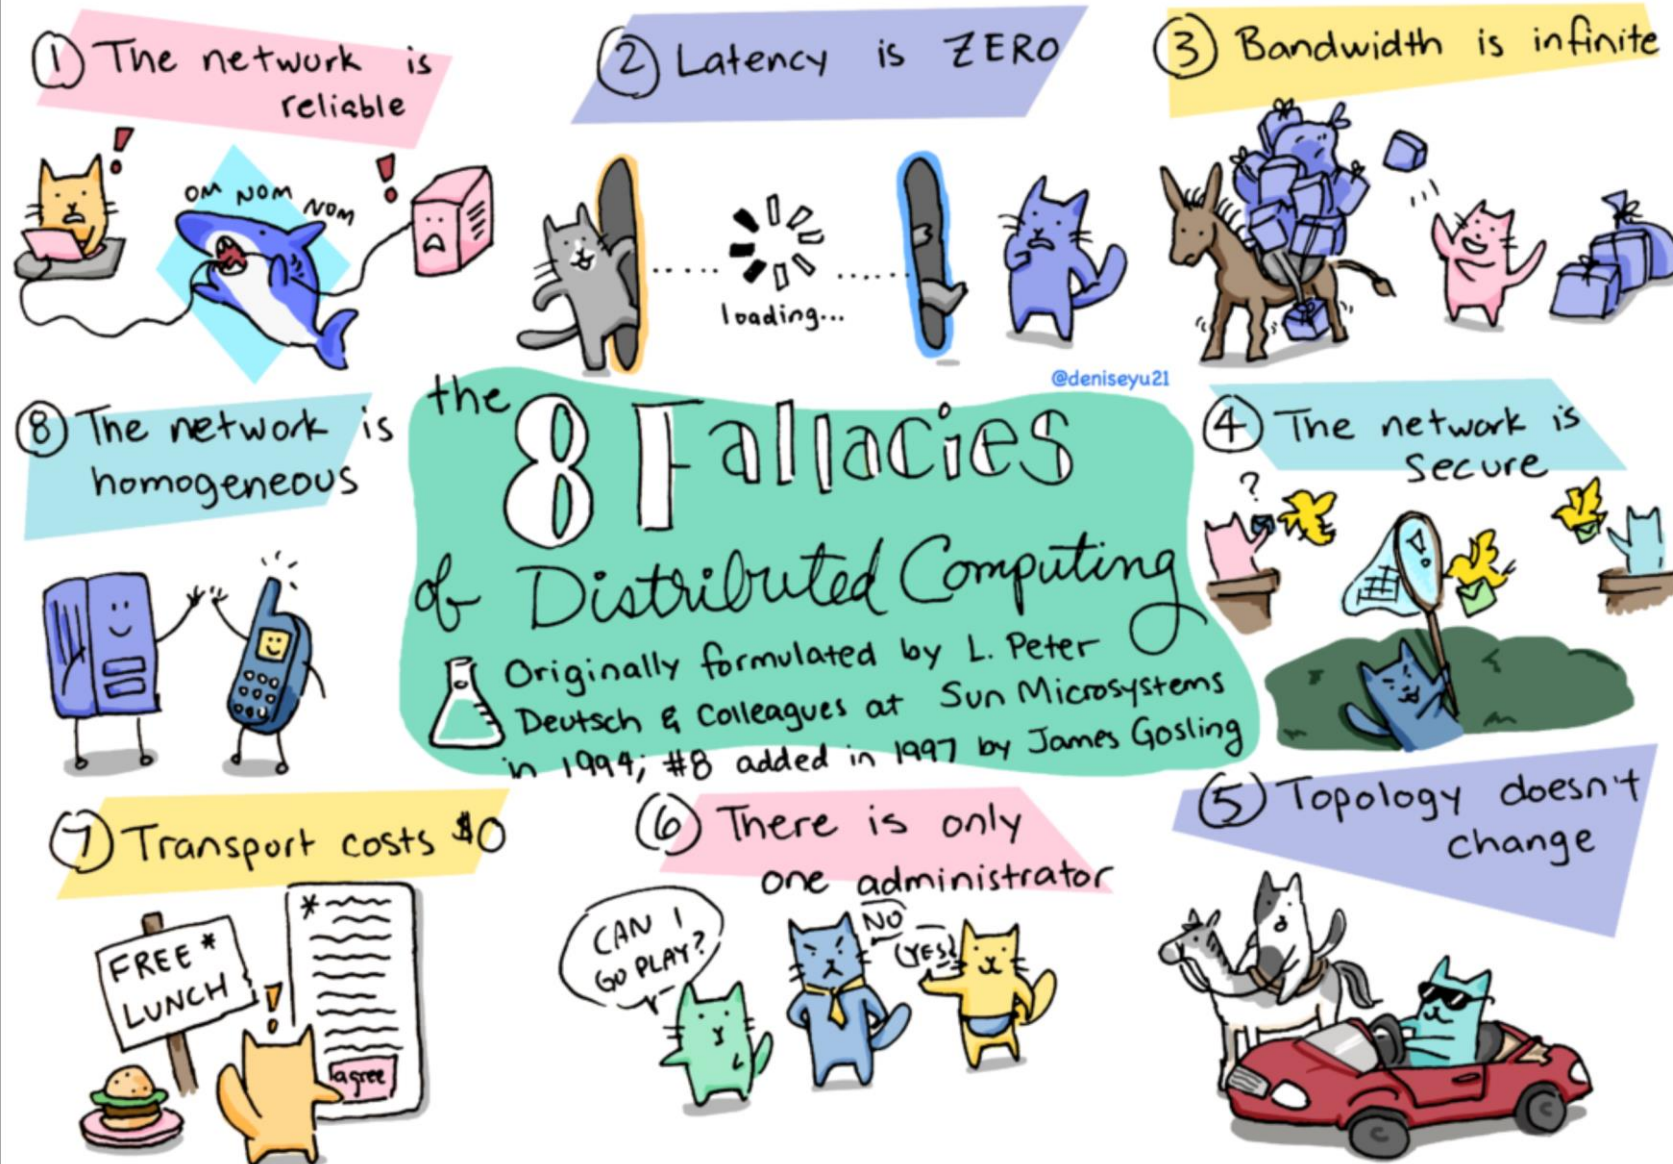
- A network that stores data on more than one physical or virtual machines at the same time
- NoSQL databases are often distributed, and the data is stored on multiple computers




# 8 Fallacies of Distributed Systems

- The 8 fallacies are
  - The network is reliable
  - Latency is zero
  - Bandwidth is infinite
  - The network is secure
  - Topology doesn't change
  - There is one administrator
  - Transport cost is zero
  - The network is homogeneous





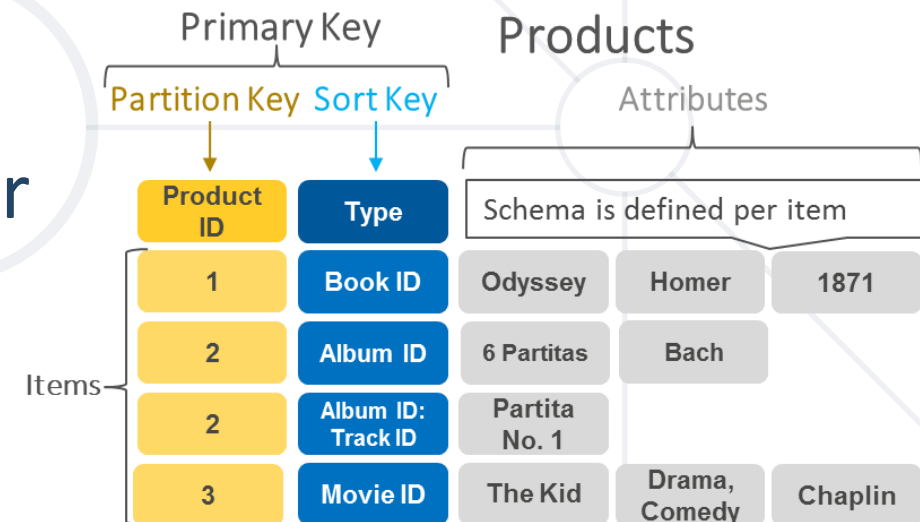


# Key-Value Databases



# Key-Value Databases


- Key-value databases work by storing and managing associative arrays
  - Keys serve as a unique identifier to retrieve an associated value
  - Values can be anything from simple objects, like integers or strings, to more complex objects






# **Document-Oriented Databases**

# Document-Oriented Databases

- 
- Document-oriented databases, or document stores, are NoSQL databases that store data in the form of documents
    - Document stores are a type of key-value store
      - Each document has a unique identifier
      - Document itself serves as the value
    - Usually stored as JSON, XML, Proto-Buff, etc.

```
{  
  "FirstName": "Bob",  
  "Address": "5 Oak St.",  
  "Hobby": "sailing"  
}
```

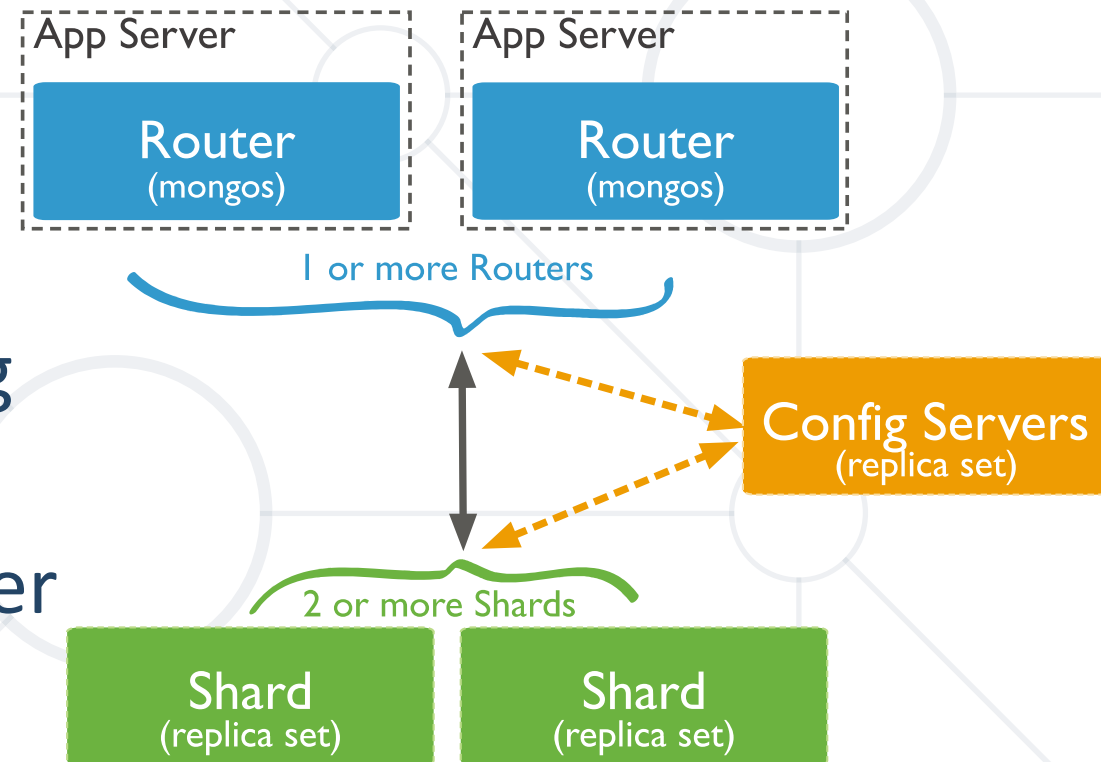
- A **replica** set is a group of mongod instances that maintain the same data set
- If the primary is unavailable, an eligible secondary will hold an election to elect itself the new primary
- All reads and writes happen from the primary (configurable)



- **Sharding** is a method for distributing data across multiple machines
- MongoDB uses sharding to support deployments with very large data sets and high throughput operations
- Database systems with large data sets or high throughput applications can challenge the capacity of a single server

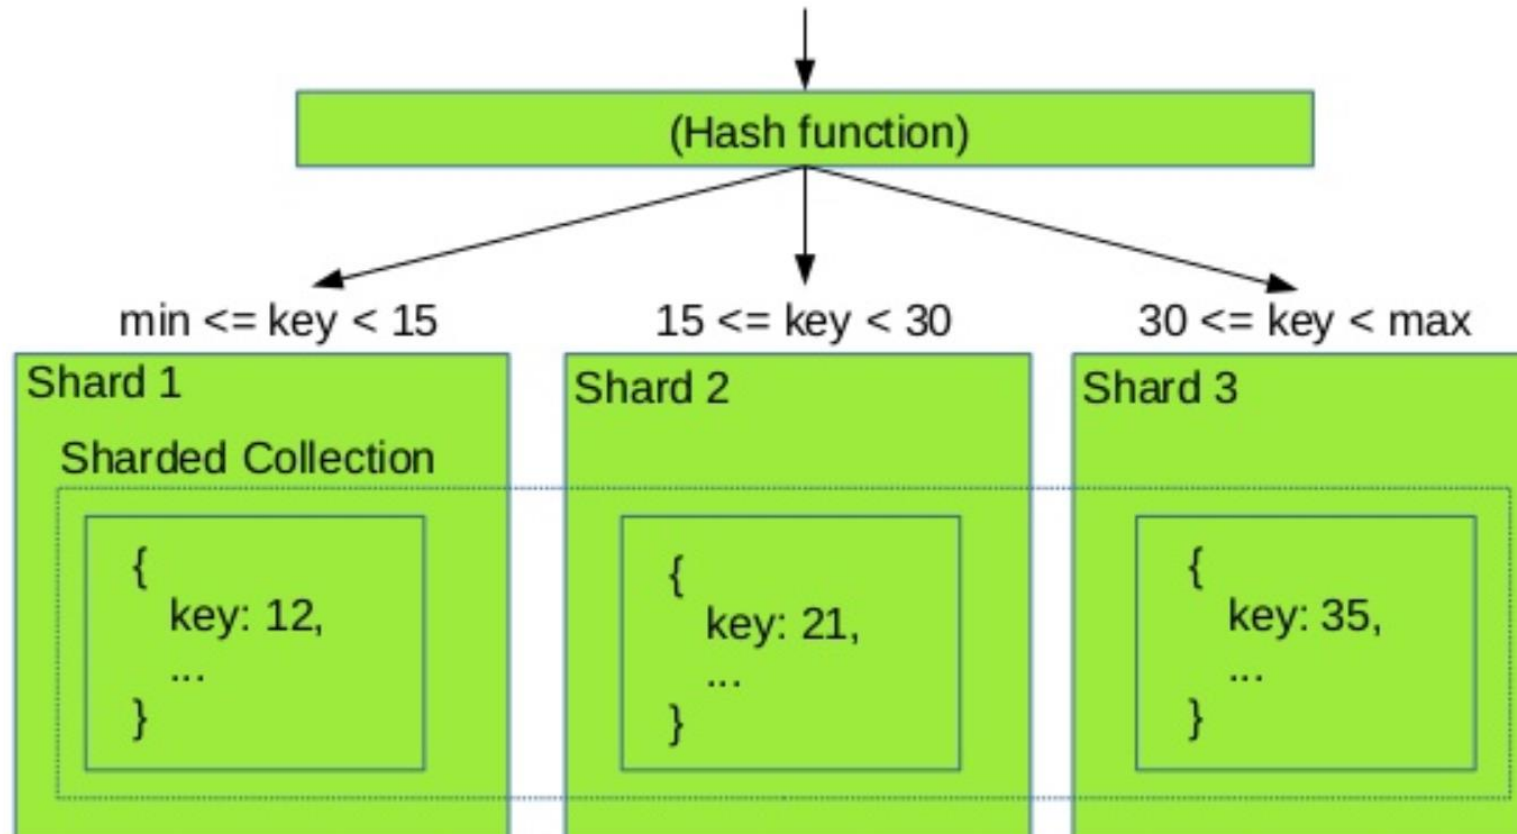
# Sharded Cluster

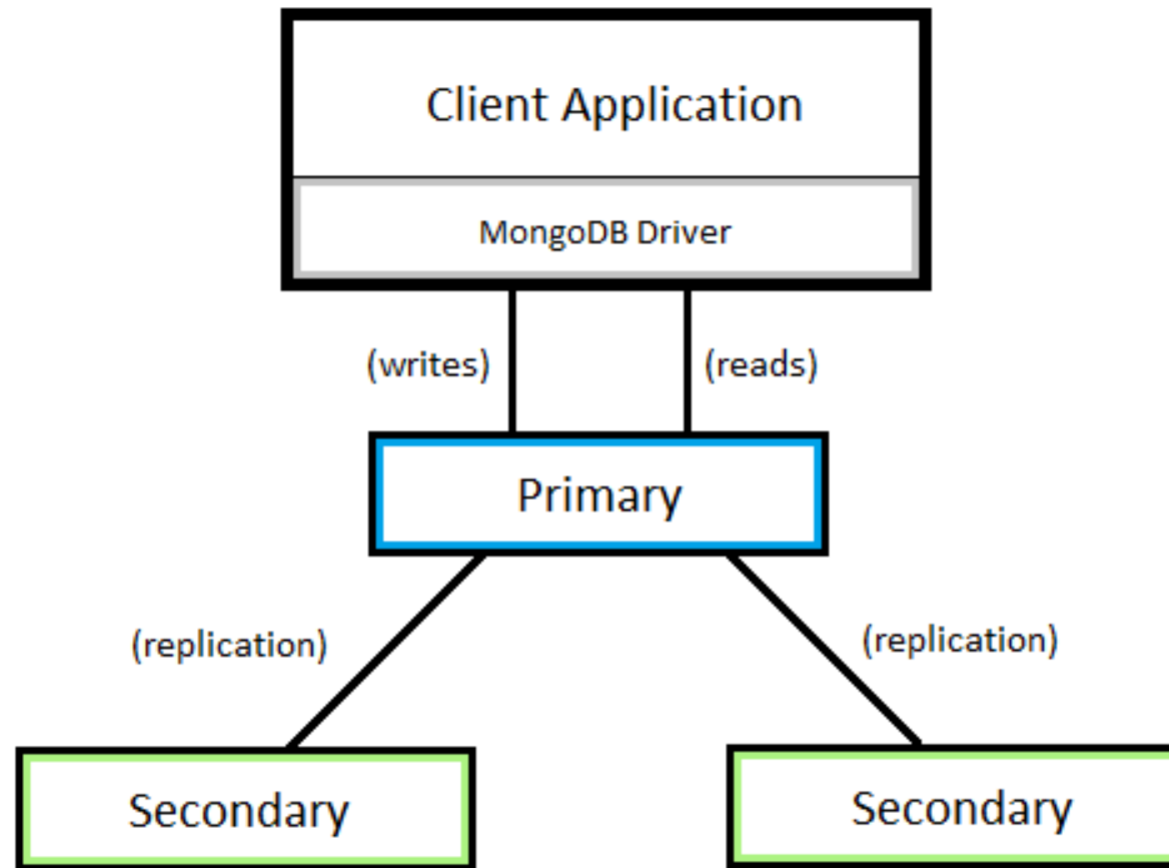
- A MongoDB **sharded cluster** consists of the following components
  - **Shard** – each contains a subset of the sharded data
  - **Mongos** - a query router, providing an interface between client applications and the sharded cluster
  - **Config Servers** - store metadata and configuration settings for the cluster



## MongoDB

Shard key selection









# Columnar Databases

# Columnar Databases (1)

- Columnar databases, are database systems that store data in columns
  - Each column is stored in a separate file or region in the system's storage
  - Examples of columnar databases are Cassandra, Hbase, Redshift, etc.

Row-oriented (1)

name	age	sex	zipcode
thomas	18	male	1416
martin	33	male	1645
bob	25	male	1613

Column-oriented (2)

name	age	sex	zipcode
thomas	18	male	1416
martin	33	male	1645
bob	25	male	1613

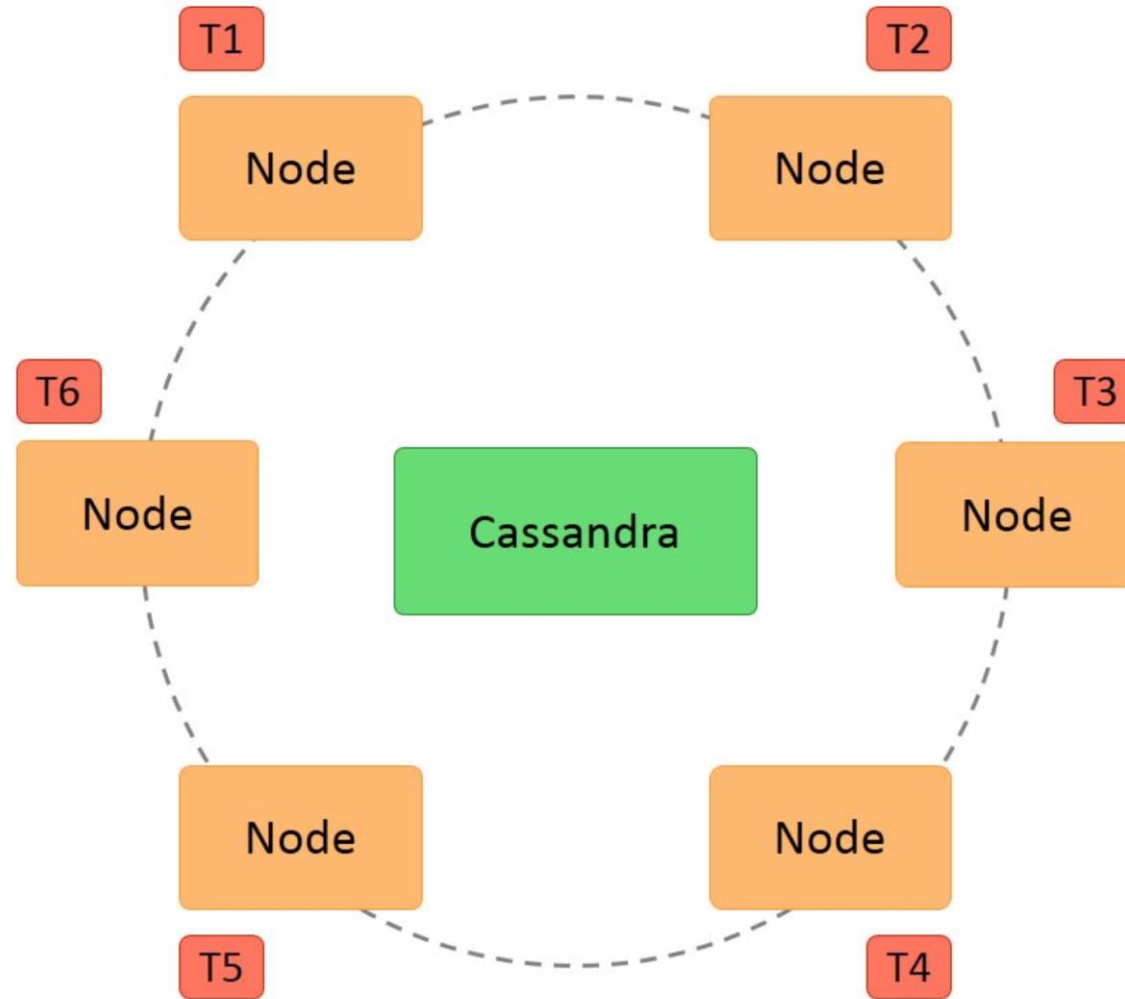
- Key benefits of column store databases include faster performance in load, search, and aggregate functions
- Column-oriented organizations are more efficient when an aggregate needs to be computed over many rows but only for a notably smaller subset of all columns of data
- Not efficient when many columns of a row are required at the same time

# Columnar Databases (3)


AuthorProfile			
Mahesh	Gender	Expertise	Rank
	Male	ADO.NET, C#, GDI+	102
	03182019	03182019	03182019
David	Gender	Book	
	Male	AWS Developer's Guide	
	03202019	03202019	
Allen	City	Book	Rank
	London	Azure Quickstart	89
	04201019	04201019	04201019

- Generally considered AP (in CAP)
- Every node in the cluster has the same role. There is no single point of failure.
- Data is distributed across the cluster (so each node contains different data)
- Failed nodes can be replaced with no downtime.
- Eventually consistent (configurable)
- Uses CQL for queries

# Cassandra (2)



- **Discord** switched to Cassandra to store billions of messages from MongoDB in November, 2015
- **Netflix** uses Cassandra as their back-end database for their streaming services
- **Apple** uses 100,000 Cassandra nodes
- **Uber** uses Cassandra to store around 10,000 features
- Many more applications



**Graph Database**




# What is a Graph Database? (1)

- Allow simple and fast retrieval of complex hierarchical structures that are difficult to model in relational systems
- No universal query language is present for graph databases (like SQL). Each database has own implementation of queries



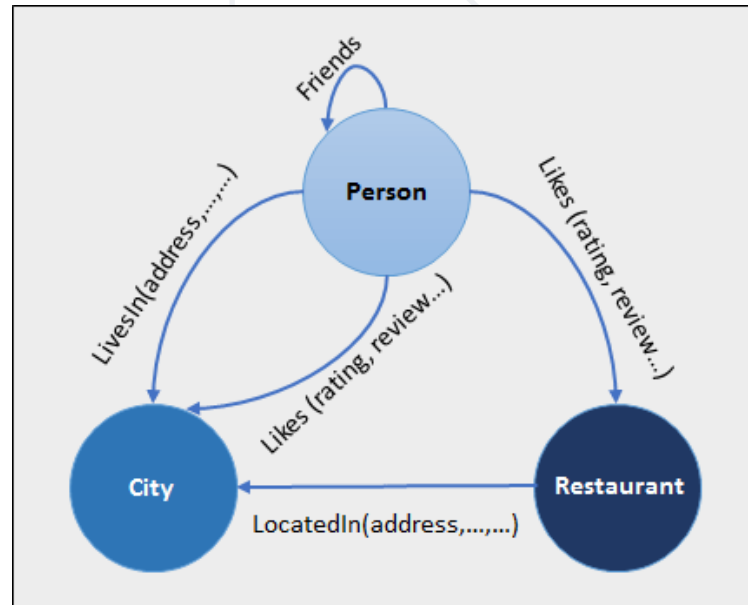
# What is a Graph Database? (2)

- A graph database contains a collection of nodes and edges
  - A node represents an object
  - An edge represents the connection or relationship between two objects



# What is a Graph Database? (3)

- Each node is identified by a unique identifier that expresses key-value pairs
- Each edge is defined by a unique identifier that details a starting or ending node, along with a set of properties






# MongoDB Overview

Installation, Configuration, Startup

# What is MongoDB?

- 
- MongoDB is a **document database**
  - It stores data in flexible, **BSON** documents
  - The document model maps to the objects in the application code, making data easy to work with
  - MongoDB is a **distributed database** at its core

```
1  {  
2    _id: "5cf0029cafff5056591b0ce7d",  
3    firstname: 'Jane',  
4    lastname: 'Wu',  
5    address: {  
6      street: '1 Circle Rd',  
7      city: 'Los Angeles',  
8      state: 'CA',  
9      zip: '90404'  
10   }  
11  }
```



- Download from: <https://www.mongodb.com/download-center>
- When **installed**, MongoDB needs a **driver**
  - One to use with Node.js, .NET, Java, etc..
  - MongoDB C#/.NET driver:  
<https://docs.mongodb.com/drivers/csharp>

- Choose one of the many
- For example
  - Robo 3T → <https://robomongo.org/download>
  - NoSQLBooster → <https://nosqlbooster.com>
  - Compass → <https://www.mongodb.com/products/compass>

# Working with MongoDB Shell Client

- Start the shell from a CLI
  - Type the command **mongo**

```
show dbs
```

Shows **all** databases  
in the data **folder**

```
use mytestdb
```

```
db.mycollection.insert({"name": "George"})
```


```
db.mycollection.find({"name": " George"})
```

```
db.mycollection.find({})
```

Gets **all** entries in  
the database

- Additional information at  
<https://docs.mongodb.com/manual/reference/mongo-shell/>





# CRUD Operations

- To **connect** to a MongoDB cluster, use the connection string for your cluster

```
using MongoDB.Bson;
using MongoDB.Driver;
...
var client = new MongoClient(
    "mongodb+srv://<username>:<password>@<cluster-address>/test?w=majority"
);
var database = client.GetDatabase("Example");
var collection = database.GetCollection<Interactions>("Interactions");
```

- To **select** a document use Linq

```
var result = IMongoCollectionExtensions  
    .AsQueryable(collection)  
    .FirstOrDefault(s => s.SiteName == "Example");
```

- **FindOneAndUpdate()**

```
var update = MongoDB.Driver.Builders.Update.Set(s => s.SiteName,  
    "New Example");  
  
collection.FindOneAndUpdate(s => s.SiteName == "Example",  
    update);
```

- **DeleteOne()**

- Deletes the first document that meets the filter

```
collection.DeleteOne(e => e.Name == "Example");
```

- **InsertOne()**

- Inserts a new document

```
collection.InsertOne(newItem);
```

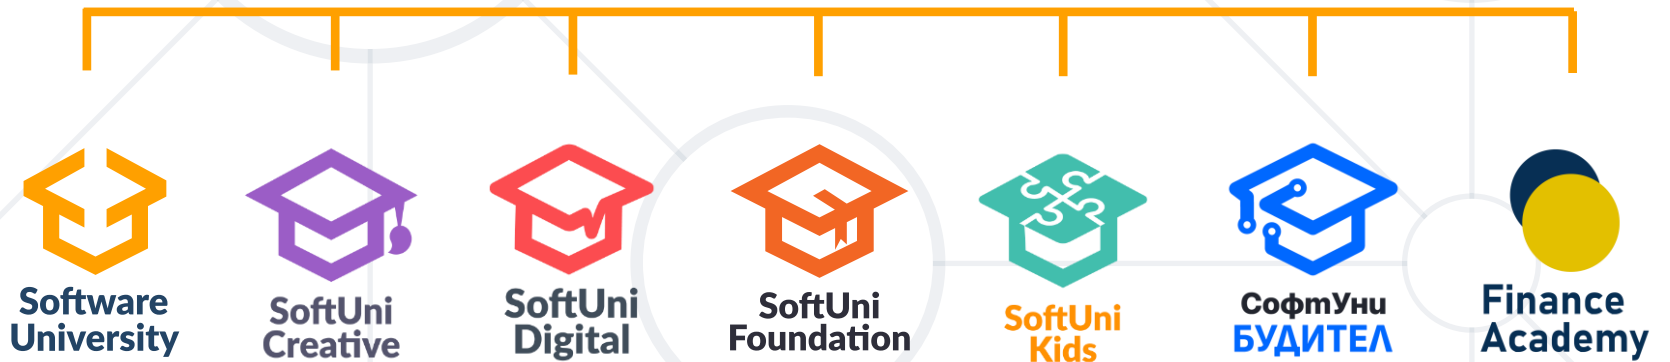
- **NoSQL**
- Ability of a system's DB to **scale up** or **down**
- **CAP** Theorem
- **Distributed** Systems
- **Key-value, Document-oriented, Columnar** and **Graph** DBs
- **MongoDB** Overview
- **CRUD** Operations



# Questions?



SoftUni



# SoftUni Diamond Partners





- Software University – High-Quality Education, Profession and Job for Software Developers

- [softuni.bg](http://softuni.bg), [about.softuni.bg](http://about.softuni.bg)

- Software University Foundation

- [softuni.foundation](http://softuni.foundation)

- Software University @ Facebook

- [facebook.com/SoftwareUniversity](https://facebook.com/SoftwareUniversity)

- Software University Forums

- [forum.softuni.bg](http://forum.softuni.bg)



- This course (slides, examples, demos, exercises, homework, documents, videos and other assets) is **copyrighted content**
- Unauthorized copy, reproduction or use is illegal
- © SoftUni – <https://about.softuni.bg>
- © Software University – <https://softuni.bg>

